

SOC542 Statistical Methods in Sociology II

Dummy, Categorical, and Non-Linear Variables

Thomas Davidson

Rutgers University

February 24, 2025

Homework

Homework 2 released after class

- ▶ Multivariate regression: specification, estimation, and interpretation
- ▶ Due Friday (2/28) at 5pm via Github Classroom

Paper proposals

- ▶ Introduction and relevant literature (~1-2 pages)
- ▶ Research question (~1 page)
 - ▶ Define your theoretical estimand
- ▶ Data (~1-2 pages)
 - ▶ Information about dataset(s)
 - ▶ Dependent variable & independent variable
 - ▶ Controls and DAG
- ▶ Methodology (~1 page)
 - ▶ Define your empirical estimand
 - ▶ Regression equation (main specification)
- ▶ References
- ▶ Due next Friday, 3/7 at 5pm (via email)

Plan

- ▶ Beyond continuous predictors
 - ▶ Dummy variables
 - ▶ Categorical variables
 - ▶ Logarithms
 - ▶ Polynomials

Dummy variables

Definitions

- ▶ A **dummy variable** is used to measure the difference between two possible states.
- ▶ Dummy variables are binary, taking a value of either zero or one.
- ▶ These values stand in for social categories of interest:
 - ▶ e.g. employed/unemployed, liberal/conservative, profit/loss

Dummy variables

Dummy variables as random variables

- ▶ We can generate dummy variables using the Binomial distribution, where $P(x = 1) = p$ and $P(x = 0) = 1 - p$.

$$x \sim \text{Binomial}(N, p)$$

Dummy variables

A simple model

```
N <- 1E5
x <- rbinom(N, 1, .4) #  $p = .4$ 
y <- 3*x + rnorm(N, 10, 1)
m <- lm(y ~ x)
round(m$coefficients,2)
```

```
## (Intercept)          x
##          10.01        2.98
```

Dummy variables

Interpretation

```
as.data.frame(cbind(x,y)) %>% group_by(x) %>%  
  summarize(mean = mean(y)) %>% as.matrix() %>% round()
```

```
##      x mean  
## [1,] 0   10  
## [2,] 1   13
```


Dummy variables

Interpretation

- ▶ The coefficient represents the expected difference in the outcome when $x = 1$ compared to $x = 0$.
- ▶ Consider the following population model, predicting income as a function of union membership.

$$\text{Income} = \beta_0 + \beta_1 \text{Union} + u$$

- ▶ β_1 represents the expected difference in income for union members compared to non-union members.
- ▶ The dummy variable also affects the interpretation of the intercept β_0 , which is the average income for non-unionized workers.

Dummy variables

Example: Union wage returns

```
gss <- haven::read_dta("../..//2022/labs/lab-data/GSS2018.dta")

data <- gss %>% select(realrinc, union, sex) %>%
  drop_na(realrinc, union) %>%
  mutate(union_dummy = ifelse(union == 1, "U", "nU"),
         sex = ifelse(sex == 1, "Male", "Female"))

u.reg <- lm(realrinc ~ union_dummy, data = data)
print(u.reg)

##
## Call:
## lm(formula = realrinc ~ union_dummy, data = data)
##
## Coefficients:
## (Intercept)  union_dummyU
##          24572          5646
```

Dummy variables

Example: Union wage returns

```
means <- data %>% group_by(union_dummy) %>%  
  summarize(m = mean(realrinc))  
print(means)
```

```
## # A tibble: 2 x 2  
##   union_dummy      m  
##   <chr>          <dbl>  
## 1 U             30218.  
## 2 nU            24572.
```

```
mean.nonunionized <- means %>% filter(union_dummy == "nU")  
print(mean.nonunionized$m + u.reg$coefficients[2])
```

```
## union_dummyU  
##           30218.2
```

Dummy variables

Reversing the reference category

```
x.rev <- ifelse(x, 0, 1)
m2 <- lm(y ~ x.rev)
round(m2$coefficients,2)

## (Intercept)      x.rev
##          13.00      -2.98
```

Dummy variables

Reference category

- ▶ The interpretation of the model depends on which value we assign to 1 or 0.
 - ▶ The value assigned to 0 is known as the **reference category**.
- ▶ For a statistical perspective, the choice is arbitrary.
- ▶ But the choices encode assumptions about the social world and can be useful for theoretical reasons.¹

¹See Johfre and Freese (2021) for further discussion of reference categories.

Dummy variables

Removing the reference category

If we estimate a model with no intercept then we get a separate parameter for each category:

$$Income = \beta_{NoUnion}^* + \beta_{Union}^* + u$$

```
u.reg.ni <- lm(realrinc ~ 0 + union_dummy, data = data)
print(coefficients(u.reg.ni))

## union_dummysU union_dummyU
##      24572.3      30218.2

diff <- coefficients(u.reg.ni)[2] - coefficients(u.reg.ni)[1]
print(diff[[1]])

## [1] 5645.907
```

Dummy variables

Removing the reference category

- ▶ In the model with no intercept we get a separate coefficient for each category.
- ▶ The difference between these coefficients is equivalent to the dummy variable in the original intercept model.

$$\beta_{Union} = \beta_{NoUnion}^* - \beta_{Union}^*$$

Dummy variables

Multiple dummy variables

- ▶ A multiple regression model can include more than one dummy variable.
- ▶ The interpretation of each coefficient is now the difference *holding other variables at their means*.
- ▶ The intercept is the mean value of the outcome when *all dummy variables are zero*.

Dummy variables

Multiple dummy variables

This model of union wage returns includes a dummy variable for sex. What is the reference category? How should we interpret the intercept?

```
u.reg2 <- lm(realrinc ~ union_dummy + sex, data = data)
print(u.reg2)

##
## Call:
## lm(formula = realrinc ~ union_dummy + sex, data = data)
##
## Coefficients:
## (Intercept)  union_dummyU      sexMale
##      20077      3699      9757
```

Categorical variables

More than two categories

- ▶ **Categorical variables** are a generalization of dummy variables to more than two categories.
 - ▶ e.g. Race/ethnicity, gender identity, highest level of education, region.
- ▶ Categories can be **ordinal**, indicating some type of numerical ranking, or **nominal**.

Categorical variables

Categorical variables as dummy variables

```
cats <- sample(c("a", "b", "c"), N, replace=TRUE,  
              prob=c(0.2, 0.2, 0.6))  
print(cats[1:10])
```

```
## [1] "c" "a" "c" "b" "a" "c" "b" "b" "c" "c"
```

```
a <- ifelse(cats == "a", 1,0)  
b <- ifelse(cats == "b", 1,0)  
c <- ifelse(cats == "c", 1,0)  
y <- 0.3*b + rnorm(N)
```

Categorical variables

Reference categories and regression results

The only difference between these models is the order. By default, the last value will be used as the reference category.

```
m4 <- lm(y ~ a + b + c)
m5 <- lm(y ~ c + a + b)
m6 <- lm(y ~ b + c + a)
```

Categorical variables

Reference categories and regression results

	(1)	(2)	(3)
(Intercept)	-0.007 (0.004)	0.305*** (0.007)	0.006 (0.007)
a	0.012 (0.008)	-0.300*** (0.010)	
b	0.312*** (0.008)		0.300*** (0.010)
c		-0.312*** (0.008)	-0.012 (0.008)
Num.Obs.	100000	100000	100000
R2	0.015	0.015	0.015
R2 Adj.	0.015	0.015	0.015
RMSE	1.00	1.00	1.00

Categorical variables

Interpretation

- ▶ Let's assume we run a survey in North America and find out respondents' country of residence. We want to estimate the following model, using USA as a reference category:

$$Income = \beta_0 + \beta_1 Canada + \beta_2 Mexico + u$$

- ▶ β_0 represents the average income for respondents in the USA.
- ▶ β_1 represents the expected difference between Canada and the USA.
- ▶ β_2 represents the expected difference between Mexico and the USA.

Categorical variables

Degrees of freedom

- ▶ Each additional category uses up a degree of freedom in our model.
 - ▶ Not a major concern unless using variables with many categories (e.g. state of residence)
- ▶ It may be defensible to treat *ordinal* variables as if they are continuous.
 - ▶ This only uses one degree of freedom.
 - ▶ Unit increases must be constant for all values and have a linear interpretation.

Categorical variables

Encoding and categorical variables

- Categorical data in surveys is often coded as numeric. This can lead to misleading results since we might consider *nominal* categories as *ordinal* (or continuous). Consider the example below using the region variable in the GSS.

```
coefficients(lm(realrinc ~ region, data = gss))
```

```
## (Intercept)      region  
## 25217.3093    -42.7065
```


Categorical variables

Encoding and categorical variables

- We can fix this by casting the variable as a factor.

```
gss$region <- as.factor(gss$region)
coefficients(lm(realrinc ~ region, data = gss))
```

```
## (Intercept)      region2      region3      region4      region5
##   29665.891   -4397.330   -5512.035   -7055.731   -3517.123
##      region7      region8      region9
##  -3834.901   -6178.623   -3194.037
```

Non-linear variables

Specifying non-linearities

- ▶ Recall that OLS regression is linear *in parameters*.
- ▶ This does not require that predictors or outcomes vary in a linear way.
- ▶ We will consider how logarithms and polynomials allow us to specify non-linear relationships between variables.

Logarithms

Logarithms and the exponential function

- ▶ The **exponential function** raises a *base* b is raised to a power² x .

$$\exp(x) = b^x$$

- ▶ The **logarithm** is the *inverse* of exponentiation.

$$\log_b(b^x) = x$$

²Note how this differs from the power function, where we raise x to a specified power. E.g. $\text{power}(x, 2) = x^2$

Logarithms

Logarithms and the exponential function

- ▶ Here are some examples of common bases:

$$\log_2(2^4) = 4$$

$$\log_{10}(10^4) = 4$$

- ▶ The **natural logarithm** uses the constant $e \approx 2.718282$ as its base:

$$\log_e(e^4) = 4$$

Logarithms

Logarithms and exponents

- ▶ We can easily verify this in R using the `log` function with specified bases:

```
log(2^4, base = 2)
```

```
## [1] 4
```

```
log(10^4, base = 10)
```

```
## [1] 4
```


```
log(exp(1)^4)
```

```
## [1] 4
```

- ▶ The default base is e . Thus, $\exp(1) = e^1 = e$.

Logarithms

Graphing logarithms




Logarithms

When to use logarithms

- ▶ Use logarithms when
 - ▶ All values of x are positive
 - ▶ x has a wide range (e.g. income)
- ▶ Interpretation
 - ▶ Logarithms allow us to transform variables to measure differences in *magnitude*
- ▶ Specification
 - ▶ Logarithms can induce normality and reduce variance for if a variable has a **log normal** distribution.


Logarithms

A unit increase of \log_{10}




Logarithms

A unit increase of \log_{10}



Logarithms


The log-normal distribution



Draws generated using `rlnorm()`.


Logarithms

The natural logarithm of the log-normal distribution




Logarithms

The distribution of 2018 GSS respondent income (realinc)



Logarithms

The distribution of 2018 GSS respondent income (realinc),
natural log



Logarithms

Logarithms and regression

- ▶ Logarithms of predictors (**Linear-log models**)
 - ▶ If we include $\log_e(x)$ as a predictor, β_i now represents the expected change associated with a unit increase in $\log_e(x)$

Logarithms

Logarithms as predictors

Here β_1 represents the effect of a one-unit increase in height *on a logarithmic scale*.

```
##  
## Call:  
## lm(formula = realrinc ~ log(height), data = gss)  
##  
## Coefficients:  
## (Intercept)    log(height)  
##      -348333          88987
```

Logarithms

Logarithms and regression

- ▶ Logarithms of outcomes (**Log-linear models**)
 - ▶ If the outcome is $\log_e(y)$. For any variable x , β_i represents the expected change in $\log_e(y)$ associated with a unit change in x .

Logarithms

Logarithms as outcomes

Here β_1 represents the effect of a one-unit increase in height on the *logarithm* of income. Note the difference in the coefficient compared to the previous model.

```
##  
## Call:  
## lm(formula = log(realrinc) ~ height, data = gss)  
##  
## Coefficients:  
## (Intercept)      height  
##      5.60849      0.06041
```

Logarithms

Logarithms and regression

- ▶ Logarithms of predictors *and* outcomes (**Log-log models**)
 - ▶ If both x and y are entered into the model as logarithms, β_i represents the expected change in $\log_e(y)$ associated with a unit change in $\log_e(x)$
 - ▶ Equivalently, this corresponds to the expected percentage change in y as a result of a 1% change in x . Hence, such coefficients can be interpreted as **elasticities**.

Logarithms

Log-log models

Here β_1 represents the effect of a one-unit increase in *logarithm* of height on the *logarithm* of income. A 1% increase in height is associated with a 4% increase in income.

```
##  
## Call:  
## lm(formula = log(realrinc) ~ log(height), data = gss)  
##  
## Coefficients:  
## (Intercept)  log(height)  
##          -7.341          4.044
```

Logarithms

Log-log models

We can still incorporate other variables into these models. Here the coefficient for sex can be interpreted as the *difference in the expected logarithm of income* between male and female respondents.

```
##  
## Call:  
## lm(formula = log(realrinc) ~ log(height) + sex, data = gss)  
##  
## Coefficients:  
## (Intercept)  log(height)          sex  
##      -0.4834      2.5136      -0.2774
```

Logarithms

Model comparison

	Log x	Log y	Log-log	Log-log+
(Intercept)	-348333.267*** (55342.266)	5.608*** (0.520)	-7.341*** (2.181)	-0.483 (3.046)
log(height)	88986.528*** (13158.504)		4.044*** (0.519)	2.514*** (0.703)
height		0.060*** (0.008)		
sex				-0.277** (0.086)
Num.Obs.	1194	1194	1194	1194
R2	0.037	0.049	0.049	0.057
R2 Adj.	0.036	0.048	0.048	0.055

+ $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Polynomials

Definitions

- ▶ **Polynomial** regression expresses non-linear relationships between continuous predictors in a linear model by adding exponents of x .
- ▶ The expected value of y is expressed as an k^{th} degree polynomial:


$$y = \beta_0 + \beta_1x + \beta_2x^2 + \beta_3x^3 + \dots + \beta_kx^k + u$$

- ▶ Generally, we use a restricted form, such as the quadratic model:

$$y = \beta_0 + \beta_1x + \beta_2x^2 + u$$

Polynomials

Quadratic functions and parabolas



$$y = 2 + -3x + 2x^2.$$

Polynomials

When to use polynomial regression


- ▶ We add polynomials to capture non-linear relationships. For example, we might expect a non-linear relationship between age and income. We could express this using the following model:

$$Income = \beta_0 + \beta_1 Age + \beta_2 Age^2 + u$$

- ▶ The effect of age is now decomposed into two coefficients:
 - ▶ β_1 captures the linear relationship between age and income.
 - ▶ β_2 captures a non-linear association between age and income.
- ▶ The coefficients no longer have a simple interpretation.
 - ▶ Can change Age while holding Age^2 constant?


Polynomials

Age and Age-Squared



Polynomials


Income and age



Age < 89 and income < 1E5. OLS fitted line.

Polynomials

Income and age



Age < 89 and income < 1E5. OLS & LOESS fitted lines.


Polynomials

Income and age

	(1)	(2)	(3)	(4)
(Intercept)	8.917*** (0.108)	5.759*** (0.294)	2.958*** (0.764)	-2.453 (1.964)
age	0.014*** (0.002)	0.166*** (0.013)	0.368*** (0.053)	0.894*** (0.184)
age2		-0.002*** (0.000)	-0.006*** (0.001)	-0.024*** (0.006)
age3			0.000*** (0.000)	0.000*** (0.000)
age4				-0.000** (0.000)
Num.Obs.	1357	1357	1357	1357
R2	0.027	0.114	0.124	0.130

Polynomials

Predictions



Predicted values from OLS models. Income measured using `realrinc`. Respondents under 89 only.

Polynomials

When to use polynomial regression

- ▶ A second-order polynomial (e.g. Age^2) is sufficient to capture non-linearity.
 - ▶ In this case, there is evidence of a curvilinear relationship between age and income.³
- ▶ Higher-order polynomial terms can improve model fit and capture more complex non-linearities but *use up degrees of freedom* and become difficult to interpret.

³For an example of polynomials used in a different context, see Dokshin, Fedor A. 2016. "Whose Backyard and What's at Issue? Spatial and Ideological Dynamics of Local Opposition to Fracking in New York State, 2010 to 2013." *American Sociological Review* 81 (5): 921–48.

Next week

Interaction terms

- ▶ What is an interaction?
- ▶ Specification
- ▶ Interpretation

Lab

- ▶ Interpreting regression coefficients