

# **SOC542 Statistical Methods in Sociology II**

## **Count outcomes**

Thomas Davidson

Rutgers University

April 7, 2025

# Course updates

## Homework

- ▶ Homework 3 grades released
- ▶ Homework 4 released after class, due next Friday, 4/18
  - ▶ Count outcomes
  - ▶ Categorical and ordered outcomes (next week)

# Course updates

## Projects

- ▶ Preliminary results due 4/25
  - ▶ One or more tables or figures of descriptive statistics
  - ▶ One or more regression tables showing
    - ▶ Bivariate results
    - ▶ Multivariate results
  - ▶ Must include at least one figure showing estimates (e.g. coefficients, predictions, marginal effects)
  - ▶ Draft write up of methodology and results
- ▶ Presentations on 5/5

# Plan

- ▶ Count outcomes
- ▶ Poisson regression
- ▶ Overdispersion and negative-binomial regression
- ▶ Offsets
- ▶ Zero-inflated models

# Count outcomes

- ▶ Count outcomes are variables defined as *non-negative integers*.
  - ▶ Values must be 0 or greater.
  - ▶ Numbers must not contain any fractional component.

# Count outcomes

- ▶ In general, we obtain count variables by counting discrete events over space and time. Many social processes produce counts:
  - ▶ How many people currently live in a census tract?
  - ▶ How many siblings does someone currently have?
  - ▶ How many times has someone ever been arrested?
  - ▶ How many sexual partners reported in one year?

# Count outcomes

## Modeling counts using OLS

- ▶ We could treat counts like continuous variables and model them using OLS.
- ▶ Such a strategy might be appropriate if a count variable is normally distributed.
- ▶ But like the LPM, we might run into problems when making predictions.
  - ▶ Predictions not constrained to be positive or counts.

# Data


## Twitter and political parties in Europe

- ▶ Data from Twitter accounts of 190 political parties in 28 countries in Europe
- ▶ Includes cumulative number of tweets and engagements (likes, replies, retweets) from 2018
- ▶ Data on left-right ideology (0-10 scale) and % of parliamentary seats held



# Data

## Twitter and political parties in Europe



# Count outcome

## Modeling counts using OLS

	OLS	OLS FE	OLS FE (Log)
Seats %	8.826 (23.091)	47.643* (21.346)	0.016 (0.009)
Ideology [0-10]	-79.080 (130.196)	-46.490 (112.614)	-0.070 (0.046)
Intercept	3066.667*** (740.226)	1958.652 (2033.170)	7.767*** (0.823)
Num.Obs.	190	190	190
R2	0.002	0.427	0.428
R2 Adj.	-0.008	0.323	0.324
F	0.228	4.112	4.120


\*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

Country FE omitted.

# Count outcome

## Making predictions with OLS


**Histogram of `predict(ols.fe, data)`**



# Count outcome

## Making predictions with OLS

**Histogram of `exp(predict(ols.fe.log, data`**



# Count outcomes

## Analyzing predictions

```
min(predict(ols.fe, data))
```

```
## [1] -370.4861
```

```
min(exp(predict(ols.fe.log, data)))
```

```
## [1] 26
```

# Poisson regression

## Modeling counts as Poisson processes

- ▶ The **Poisson** distribution is a discrete probability distribution that indicates the count of events in a fixed interval. These counts can be considered as rates of events per unit.<sup>1</sup>
- ▶ The *probability mass function* is defined by a single parameter  $\lambda$ , where the probability of observing  $k$  events is equal to

$$P(x = k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

- ▶ For any Poisson distributed random variable,  $x$

$$E(x) = \lambda = \text{Var}(x)$$

---

<sup>1</sup>The distribution gets its name from French mathematician Siméon Denis Poisson [1781-1840].

# Poisson regression

## Modeling counts as Poisson processes

- ▶ Let's say the average number of visits to doctor each year is 1.6.
- ▶ We can model the probabilities of observing different numbers of visits given  $\lambda = 1.6$ :

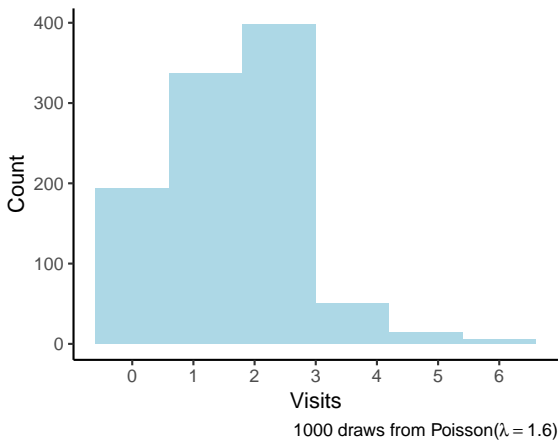
$$P(k \text{ visits a year}) = \frac{1.6^k e^{-1.6}}{k!}$$

$$P(0 \text{ visits a year}) = \frac{1.6^0 e^{-1.6}}{0!} = \frac{e^{-1.6}}{1} \approx 0.2$$

$$P(1 \text{ visits a year}) = \frac{1.6^1 e^{-1.6}}{1!} = \frac{1.6 e^{-1.6}}{1} \approx 0.4$$

# Poisson regression

## Poisson distributions





# Poisson regression

Poisson distributions,  $E[x] = \lambda = \text{Var}(x)$

```
round(mean(x),2)
```

```
## [1] 1.59
```

```
round(var(x),2)
```

```
## [1] 1.53
```

## Poisson regression

- ▶ The Poisson regression model assumes that the outcome is Poisson distributed, conditional on the observed predictors.

$$y \sim \text{Poisson}(\lambda)$$

- ▶ To ensure that our estimates are positive, we can use a logarithmic *link function*, thus

$$y = \log(\lambda) = \beta_0 + \beta_1 x_1 + \beta_2 x_1 + \dots + \beta_k x_k$$

- ▶ Like logistic regression, this equation can equivalently be expressed using the *inverse* of the logarithm function:

$$\lambda = e^{\beta_0 + \beta_1 x_1 + \beta_2 x_1 + \dots + \beta_k x_k}$$

# Poisson regression

## Fitting a model

```
pois <- glm(tweet_count ~ seats_per + left_right +  
            country,  
            data = data, family = poisson(link = "log"))
```

## Poisson regression

	OLS FE (Log)	Poisson
Seats %	0.016 (0.009)	0.013*** (0.000)
Ideology [0-10]	-0.070 (0.046)	-0.018*** (0.001)
Intercept	7.767*** (0.823)	7.708*** (0.012)
Num.Obs.	190	190
R2	0.428	
R2 Adj.	0.324	
Log.Lik.	-308.005	-166974.776

\*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

Country FE omitted.

# Poisson regression

## Interpretation

- ▶ The intercept  $\beta_0$  is the *logged* average value of the outcome when all other predictors are equal to zero.
- ▶ Each coefficient  $\beta_i$  indicates the effect of a unit change of  $x_i$  on the *logarithm* of the outcome.
  - ▶ e.g.,  $\beta_{seats\%} = 0.013$  implies that the expected log number of tweets increases by 0.02 in response to a 1-unit, or 1% increase in parliamentary seats held by a party.
- ▶ Coefficients can be interpreted as *multiplicative* changes after exponentiation
  - ▶ e.g.,  $e^{\beta_{seats\%}} = e^{0.013} \approx 1.013$ . This implies that a ~1.3% increase in tweets.
  - ▶ These coefficients are sometimes referred to as **incident rate ratios (IRRs)**.

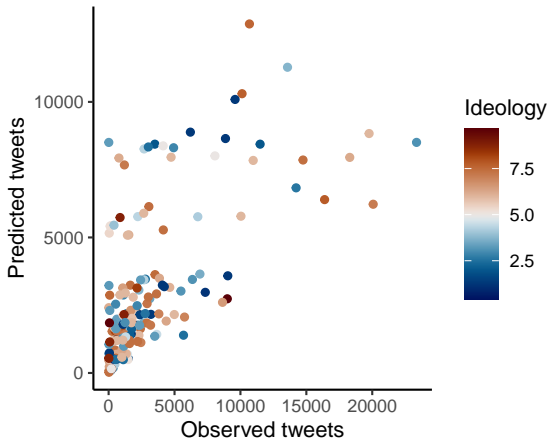
## Poisson regression

	Poisson (Exponentiated)
Seats %	1.014*** (0.000)
Ideology [0-10]	0.983*** (0.001)
Intercept	2225.822*** (25.924)
Num.Obs.	190
Log.Lik.	-166974.776
F	11552.594


\*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

Country FE omitted.

# Poisson regression

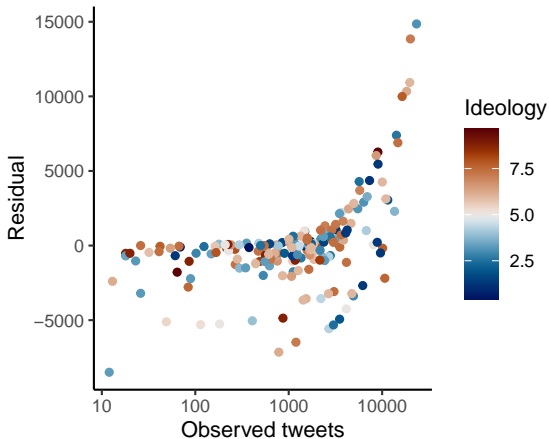


# Poisson regression

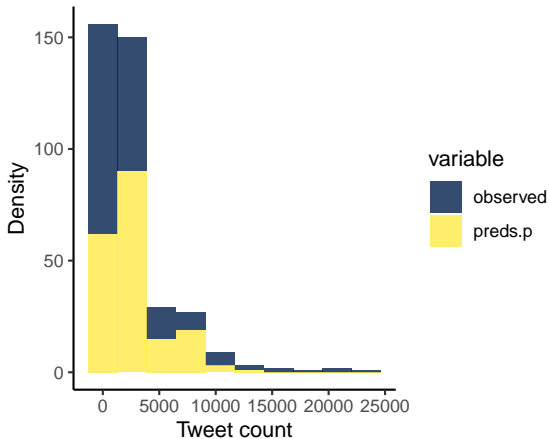




# Poisson regression




# Poisson regression



# Overdispersion

- ▶ A random variable is **overdispersed** if the observed variability is greater than the variability expected by the underlying probability model.
- ▶ **Underdispersion** occurs if the variability is lower than expected, but it is rarely an issue.

# Overdispersion



Poisson( $\lambda = 1.6$ ) and NegBin( $\mu = 1.6, \theta = 0.5$ )

# Overdispersion

## Negative binomial distribution and regression

- ▶ The **negative binomial** distribution (aka the gamma-Poisson distribution) includes an additional parameter  $\theta$  to account for dispersion, referred to as a **scale parameter**.

$$y = \text{NegativeBinomial}(\lambda, \theta)$$

- ▶ In negative binomial regression,  $\theta$  is estimated from the data. The value must be positive.
  - ▶ Lower values indicate greater overdispersion.
  - ▶ Negative binomial becomes identical to Poisson as  $\lim_{\theta \rightarrow \infty}$ .

# Overdispersion

## Fitting a negative binomial regression

Negative binomial regression is not implemented in `glm`. Instead, we can use the `glm.nb` function from the MASS package.<sup>2</sup>

```
library(MASS)
nb <- glm.nb(tweet_count ~ seats_per + left_right + country,
              data = data)
```

---

<sup>2</sup>The “fixest” package has an implementation, “fenegbin” that is more suitable for these data as it can also cluster standard errors.

## Comparing Poisson and negative binomial regression

	Poisson	Negative binomial
Seats %	0.013*** (0.000)	0.014* (0.006)
Ideology [0-10]	-0.018*** (0.001)	-0.054 (0.031)
Intercept	7.708*** (0.012)	8.031*** (0.564)
Num.Obs.	190	190
Log.Lik.	-166974.776	-1602.681
F	11552.594	9.208

\*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

Country FE omitted. Exponentiated coefficients.

# Negative binomial regression

```
nb$theta
```

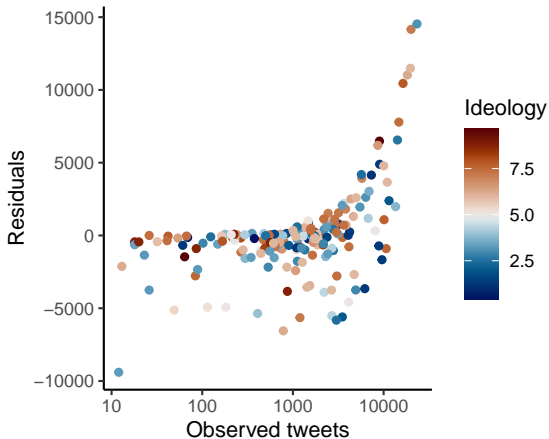
```
## [1] 1.197566
```

```
nb$SE.theta
```

```
## [1] 0.110282
```



# Negative binomial regression




# Negative binomial regression

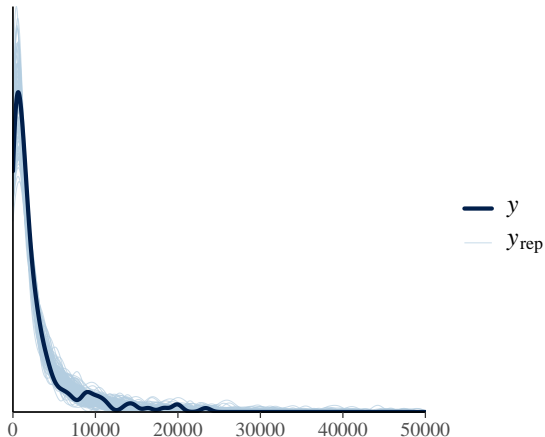
## Bayesian estimation

```
pois.b <- stan_glm(tweet_count ~ seats_per + left_right +  
                  country,  
                  data = data,  
                  family = poisson,  
                  seed = 08901, chains = 1, iter = 4000, refresh = 0)  
  
nb.b <- stan_glm(tweet_count ~ seats_per + left_right +  
                country,  
                data = data,  
                family = neg_binomial_2(),  
                seed = 08901, chains = 1, iter = 4000, refresh = 0)
```


# Poisson posterior predictive check




# Negative binomial posterior predictive check



# Poisson PSIS plot



# Negative binomial PSIS plot



# Negative binomial regression


## Comparing Poisson and negative binomial models

```
loo_compare(l.pois, l.nb)
```

```
##           elpd_diff se_diff  
## nb.b           0.0      0.0  
## pois.b -170636.7    22413.8
```

# Negative binomial regression

## Bayesian estimate of $\theta$





# Offsets

## Intuition

- ▶ Assume a count outcome  $y$  is measured over varying time intervals  $t$ . The level of  $y$  will vary both as a function of the underlying count process and the length of **exposure**.<sup>3</sup>
- ▶ We can add an **offset** to our model to account for varying exposures.
- ▶ The outcome of a model with an offset is now  $\frac{y}{t}$ .

---

<sup>3</sup>The same logic would apply if we measured quantities over varying spatial units, e.g. counting people in blocks versus census tracts.

# Offsets

## Explanation

- ▶ The mean of a Poisson process,  $\lambda$  is implicitly  $\lambda = \frac{\mu}{\tau}$ , the expected number of events,  $\mu$ , over the duration  $\tau$ .
- ▶ Assume a Poisson process where  $\lambda_i$  is the expected number of events for the  $i^{th}$  observation. We can write the link function as

$$y = \text{Poisson}(\lambda)$$

$$\log(\lambda) = \log\left(\frac{\mu}{\tau}\right) = \beta_0 + \beta_1 x$$

- ▶ This can be re-written as

$$= \log(\mu) - \log(\tau) = \beta_0 + \beta_1 x$$

# Offsets

## Explanation

- ▶ We can think of  $\tau$  as the number of **exposures** for each observation. Thus, we can write out a new model for  $\mu$ :

$$y \sim \text{Poisson}(\mu)$$

$$\log(\mu) = \log(\tau) + \beta_0 + \beta_1 x$$

# Offsets

## Example: Predicting retweet rates

- ▶ The number of retweets depends on the number of times a party tweeted
  - ▶ No tweets, no retweets
  - ▶ More tweets, more retweets?
- ▶ Two specifications
  - ▶ No offset
  - ▶  $\text{Log}(\text{tweets})$  included as offset

# Offsets

## Example: Predicting retweet rates

```
nb.rt <- glm.nb(retweet_total ~  
                seats_per + left_right + country,  
                data = data)  
nb.rt.o <- glm.nb(retweet_total ~ offset(log(tweet_count)) +  
                  seats_per + left_right + country,  
                  data = data)
```

# Offsets

	NB	NB (Offset)
Seats %	1.033*** (0.009)	1.033*** (0.007)
Ideology [0-10]	0.973 (0.045)	0.977 (0.034)
Intercept	21286.898*** (17921.337)	7.528** (4.739)
Num.Obs.	190	190
Log.Lik.	-2164.836	-2090.518
F	21.379	13.095

\*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

# Offsets

## Using offsets

- ▶ Offsets allow models to be interpreted as rates rather than counts.
- ▶ Always include an offset if there are differences in measurement intervals across observations.
- ▶ Offsets can also be included when intervals are constant if a rate is more informative.

# Zero-inflated models

## Intuition

- ▶ Some count outcomes have high rates of zeros. What if zeros are generated by a different process compared to non-zeros?
- ▶ **Zero-inflated models** allow us to jointly model the process determining whether counts are non-zero and the expected count for each non-zero observation.



# Zero-inflated models

## Specification

- ▶ The zero-inflated Poisson model consists of a mixture of two linear models, a logistic regression predicting the probability of a zero and a Poisson model predicting the count outcome.

$$y_i = ZIPoisson(p, \lambda)$$

$$\text{logit}(p) = \gamma_0 + \gamma_1 z$$

$$\log(\lambda) = \beta_0 + \beta_1 x$$

- ▶ Each model has its own set of regression parameters. These can be specified differently to model each process.

# Zero-inflated models

## Example: Books borrowed from the library

- ▶ Are you borrowing any books from the library?
  - ▶ If so, how many?

# Zero-inflated models

## Example: Books borrowed from the library

```
N <- 1000 # N

z <- rnorm(N, 0.5, 1) # Random variable determines library usage=
p_lib <- 1/(1 + (exp(1)^-(z))) # Convert to probability

lib <- rep(0,N) # Generate binary library variable
for (i in 1:N) {
  lib[i] <- rbinom(1, 1, p_lib[i])
}

sum(lib)/N

## [1] 0.604
```

# Zero-inflated models

## Example: Books borrowed from the library

```
x <- rnorm(N) # Random variable for books borrowed

books <- c() # Store number of books borrowed for each student
for (i in 1:N) {
  if (lib[i] == 1) { # Borrow books if library visitor
    books[i] <- rpois(1, lambda = exp(0.5 + x[i]))
  } else {books[i] <- 0} # Otherwise zero
}
mean(books)

## [1] 1.584

max(books)

## [1] 52
```

# Zero-inflated models

## Two kinds of zeros

```
sum(books == 0)
```

```
## [1] 542
```


```
sum(books == 0 & lib == 1)
```

```
## [1] 146
```

```
sum(books == 0 & lib == 0)
```

```
## [1] 396
```

# Zero-inflated models

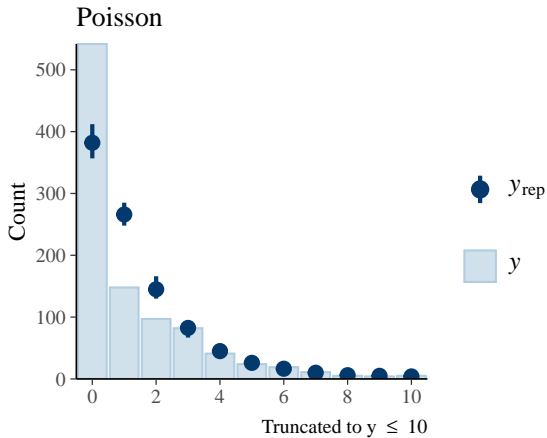


# Zero-inflated models

## Estimating a Poisson model

```
book.data <- as.data.frame(cbind(books, x, z))  
  
pois.m <- stan_glm(books ~ x, data = book.data, family = poisson(),  
  seed = 08901, chains = 1, refresh = 0)
```


# Poisson posterior predictive checks





# Poisson posterior predictive checks

Predicted proportion of zeros



# Zero-inflated models

## Estimating a zero-inflated Poisson model

We must use the `brms` library to implement Bayesian zero-inflated Poisson regression.

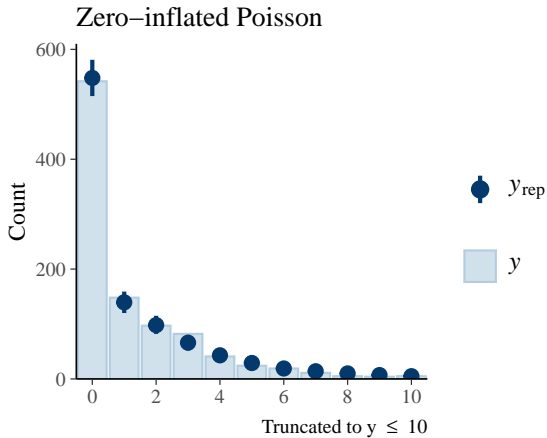
```
library(brms)
zip <- brm(bf(books ~ x,
             zi ~ z),
           data = book.data,
           family = zero_inflated_poisson(link = "log",
                                           link_zi = "logit"),
           seed = 08901, refresh = 0, chains = 1, backend = "cmdstanr")

## Running MCMC with 1 chain...
##
## Chain 1 finished in 6.0 seconds.
```


## Zero-inflated models

	Poisson	ZI Poisson
(Intercept)	-0.009 [-0.082, 0.061]	
x	0.988 [0.939, 1.039]	
b_Intercept		0.519 [0.445, 0.597]
b_x		0.992 [0.940, 1.041]
b_zi_Intercept		-0.099 [-0.283, 0.086]
b_zi_z		-0.865 [-1.059, -0.651]
Num.Obs.	1000	1000

# Posterior predictive checks



# Posterior predictive checks



# Zero-inflated models

## Comparing standard and zero-inflated Poisson models

##	elpd_diff	se_diff
## zip	0.0	0.0
## pois.m	-479.1	45.1

# Summary

- ▶ Standard linear models are generally unsuitable for count data
- ▶ Poisson regression can be used for many count outcomes
- ▶ Overdispersion occurs when variation higher than expected under Poisson model
  - ▶ Negative binomial regression includes a scale parameter to model this
- ▶ Offsets transform from counts to rates and should be used when measurement intervals vary
- ▶ Zero-inflated models can decompose processes generating zeros and counts

## Next week

- ▶ Categorical outcomes
  - ▶ Multinomial and ordered logistic regression