

# **SOC542 Statistical Methods in Sociology II**

## **Count outcomes**

Thomas Davidson

Rutgers University

April 4, 2022

# Course updates

- ▶ Homework 4 will be released new week
  - ▶ Count outcomes
  - ▶ Categorical and ordered outcomes

# Plan

- ▶ Count outcomes
- ▶ Poisson regression
- ▶ Overdispersion and negative-binomial regression
- ▶ Offsets
- ▶ Zero-inflated models

# Count outcomes

- ▶ Count outcomes are variables defined as *non-negative integers*.
  - ▶ Values must be 0 or greater.
  - ▶ Numbers must not contain any fractional component.

# Count outcomes

- ▶ In general, we obtain count variables by counting discrete events over space and time. Many social processes produce counts:
  - ▶ How many people live in a census tract?
  - ▶ How many siblings does someone have?
  - ▶ How many times has someone been arrested?

# Count outcome

## Modeling counts using OLS

- ▶ We could treat counts like continuous variables and model them using OLS.
- ▶ Such a strategy might be appropriate if a count variable is normally distributed.
  - ▶ This could occur if a continuous variable was rounded.
- ▶ But like the LPM, we might run into problems when making predictions:

# Data


## Twitter and political parties in Europe

```
data <- read_csv("data/twitter_parties_2016.csv") #  
data <- data %>%  
  replace_na(list(left_right = 5,  
                  seats_per = 0))
```

# Data

## Twitter and political parties in Europe

Histogram of data\$tweet\_count





# Count outcome


## Modeling counts using OLS

	OLS	OLS FE	OLS FE (Log)
(Intercept)	2372.846*** (555.477)	1341.047 (1267.149)	6.557*** (0.563)
populist	1457.289* (625.378)	1184.876* (538.252)	0.264 (0.239)
left_right	-79.666 (99.156)	-76.613 (86.319)	-0.051 (0.038)
seats_per	24.056 (19.043)	52.538** (16.387)	0.022** (0.007)
Num.Obs.	255	255	255
R2	0.029	0.428	0.419
R2 Adj.	0.018	0.351	0.341
Log.Lik.	-2451.473	-2384.101	-415.757
F	2.520	5.580	5.375

Country FE omitted.

# Count outcome


## Making predictions with OLS



# Count outcome

## Making predictions with OLS

Histogram of `exp(predict(ols.fe.log, data`



# Poisson regression

## Modeling counts as Poisson processes

- ▶ The **Poisson** distribution is a discrete probability distribution that indicates the number of events in a fixed time or space. These counts can be considered as rates of events per unit.<sup>1</sup>
- ▶ The *probability mass function* is defined by a single parameter  $\lambda$ , where the probability of observing  $k$  events is equal to

$$P(x = k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

- ▶ For any Poisson distributed random variable,  $x$

$$E(x) = \lambda = \text{Var}(x)$$

---

<sup>1</sup>The distribution gets its name from French mathematician Siméon Denis Poisson [1781-1840])

# Poisson regression

## Modeling counts as Poisson processes

- ▶ Let's say the average number of visits to the dentist in a single year is 1.6.
- ▶ We can model the probabilities of observing different numbers of visits given  $\lambda = 1.6$ :


$$P(k \text{ visits a year}) = \frac{1.6^k e^{-1.6}}{k!}$$

$$P(0 \text{ visits a year}) = \frac{1.6^0 e^{-1.6}}{0!} = \frac{e^{-1.6}}{1} \approx 0.2$$

$$P(1 \text{ visits a year}) = \frac{1.6^1 e^{-1.6}}{1!} = \frac{1.6 e^{-1.6}}{1} \approx 0.4$$

# Poisson regression

## Poisson distributions



1000 draws from  $\text{Poisson}(\lambda = 1.6)$

# Poisson regression

Poisson distributions,  $E[x] = \lambda = \text{Var}(x)$

```
round(mean(x),2)
```

```
## [1] 1.57
```

```
round(var(x),2)
```

```
## [1] 1.54
```

## Poisson regression

- ▶ The Poisson regression model assumes that the outcome is Poisson distributed, conditional on the observed predictions.

$$y \sim \text{Poisson}(\lambda)$$

- ▶ To ensure that our estimates are positive, we can use a logarithmic *link function*, thus

$$y = \log(\lambda) = \beta_0 + \beta_1 x_1 + \beta_2 x_1 + \dots + \beta_k x_k$$

- ▶ Like logistic regression, this equation can equivalently be expressed using the *inverse* of the logarithm function:

$$\lambda = e^{\beta_0 + \beta_1 x_1 + \beta_2 x_1 + \dots + \beta_k x_k}$$



# Poisson regression

## Fitting a model

```
pois <- glm(tweet_count ~ populist + left_right +  
            seats_per + country,  
            data = data, family = poisson(link = "log"))
```

## Poisson regression

	OLS FE (Log)	Poisson
(Intercept)	6.557*** (0.563)	7.233*** (0.010)
populist	0.264 (0.239)	0.354*** (0.003)
left_right	-0.051 (0.038)	-0.019*** (0.001)
seats_per	0.022** (0.007)	0.018*** (0.000)
Num.Obs.	255	255
R2	0.419	
R2 Adj.	0.341	
Log.Lik.	-415.757	-211788.969

Country FE omitted.

+  $p < 0.1$ , \*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

# Poisson regression

## Interpretation

- ▶ The intercept  $\beta_0$  is the *logged* average value of the outcome when all other predictors are equal to zero.
- ▶ Each coefficient  $\beta_i$  indicates the effect of a unit change of  $x_i$  on the *logarithm* of the outcome.
  - ▶ e.g.,  $\beta_{populism} = 0.354$  implies that the expected log number of tweets for populist parties is higher than non-populists by 0.354.
- ▶ Coefficients can be interpreted as *multiplicative* changes after exponentiation
  - ▶ e.g.,  $e^{\beta_{populism}} = e^{0.354} \approx 1.425$ . This implies that populist parties tweet 1.425 times as frequently or 42.5% more frequently than non-populists.
  - ▶ These coefficients are sometimes referred to as `textbf{incident rate ratios (IRRs)}`.

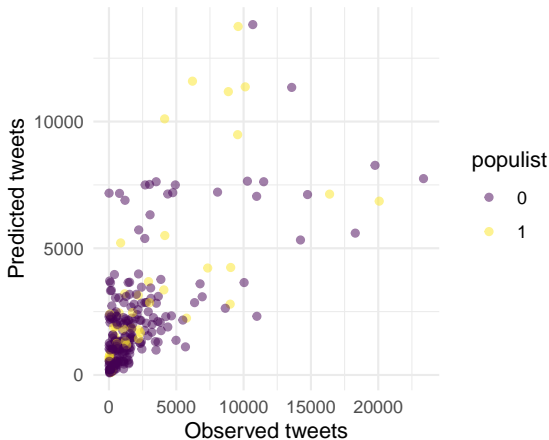
# Poisson regression

	Poisson
(Intercept)	1384.634*** (0.010)
populist	1.425*** (0.003)
left_right	0.982*** (0.001)
seats_per	1.018*** (0.000)
Num.Obs.	255
Log.Lik.	-211788.969
F	14872.576

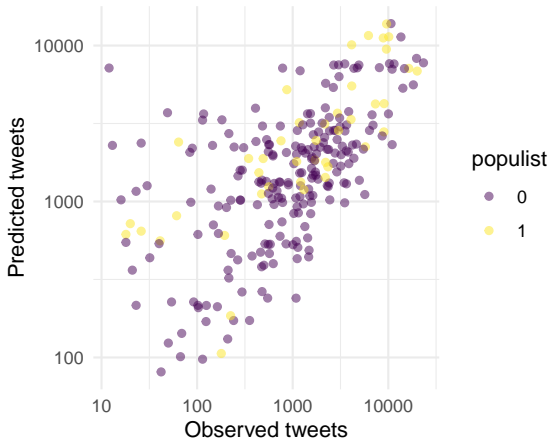
Country FE omitted.

+  $p < 0.1$ , \*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

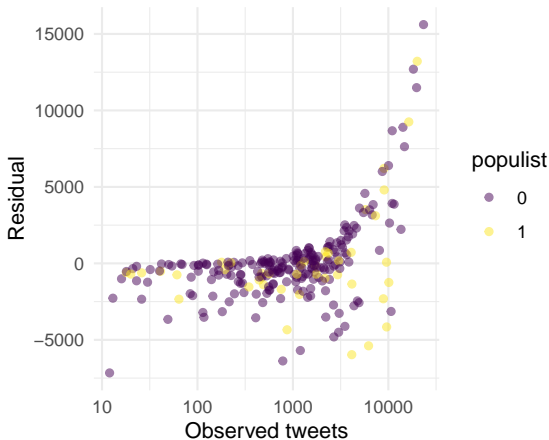
# Poisson regression




# Poisson regression



# Poisson regression



# Poisson regression

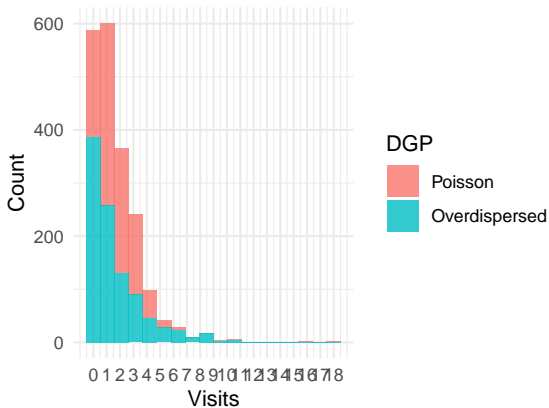




# Overdispersion

- ▶ A random variable is **overdispersed** if the observed variability is greater than the variability expected by the underlying probability model.
- ▶ In this case, we can see that the variance is far larger than the mean.
  - ▶ We could see this in the descriptive statistics, but the issue can only be properly diagnosed after fitting a model (note that the variance of the data is more than two times as large as the predicted values)
- ▶ **Underdispersion** occurs if the variability is lower than expected, but it is rarely an issue.

# Overdispersion



Poisson( $\lambda = 1.6$ ) and NegBin( $\mu = 1.6, \theta = 1$ )

# Overdispersion

## Negative binomial distribution and regression

- ▶ The **negative binomial** distribution (aka the gamma-Poisson distribution) includes an additional parameter  $\theta$  to account for dispersion, referred to as a **scale parameter**.

$$y = \text{NegativeBinomial}(\lambda, \theta)$$

- ▶ In negative binomial regression,  $\theta$  is estimated from the data. The value must be positive.
  - ▶ Lower values indicate greater overdispersion.
  - ▶ Negative binomial becomes Poisson as  $\lim_{\theta \rightarrow \infty}$ .

# Overdispersion

## Fitting a negative binomial regression

The procedure for estimating a negative binomial regression via Maximum Likelihood is not implemented in `glm`. Instead, we use the modified `glm.nb` function from the `MASS` package.

```
library(MASS)
nb <- glm.nb(tweet_count ~ populist + left_right +
              seats_per + country,
              data = data)
```

## Comparing Poisson and negative binomial regression

	Poisson	Negative binomial
(Intercept)	1384.634*** (0.010)	1614.471*** (0.410)
populist	1.425*** (0.003)	1.320 (0.174)
left_right	0.982*** (0.001)	0.950+ (0.028)
seats_per	1.018*** (0.000)	1.023*** (0.005)
Num.Obs.	255	255
Log.Lik.	-211788.969	-2120.171
F	14872.576	9.935

Country FE omitted. Exponentiated coefficients.

+  $p < 0.1$ , \*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

# Negative binomial regression


```
nb$theta
```

```
## [1] 1.0883
```

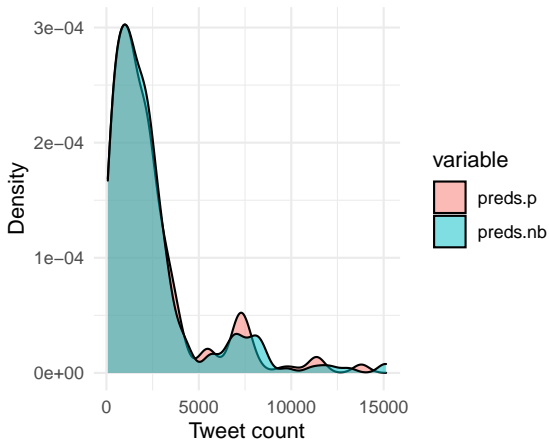
```
nb$SE.theta
```

```
## [1] 0.08578371
```

# Negative binomial regression



# Negative binomial regression






# Negative binomial regression


## Bayesian estimation

```
pois.b <- stan_glm(tweet_count ~ populist + left_right +  
                  seats_per + country,  
                  data = data,  
                  family = poisson,  
                  seed = 08901, chains = 1, refresh = 0)  
  
nb.b <- stan_glm(tweet_count ~ populist + left_right +  
                seats_per + country,  
                data = data,  
                family = neg_binomial_2(),  
                seed = 08901, chains = 1, refresh = 0)
```

# Poisson posterior predictive check




# Negative binomial posterior predictive check




# Poisson PSIS plot

PSIS diagnostic plot



# Negative binomial PSIS plot



# Negative binomial regression


## Comparing Poisson and negative binomial models

```
loo_compare(l.pois, l.nb)

##           elpd_diff se_diff
## nb.b           0.0      0.0
## pois.b -215832.9    24353.1
```

# Negative binomial regression

## Bayesian estimate of $\theta$



# Offsets

## Intuition

- ▶ Assume a count outcome  $y$  is measured over varying time intervals  $t$ . The level of  $y$  will vary both as a function of the underlying count process and the length of **exposure**.<sup>2</sup>
- ▶ We can add an **offset** to our model to account for varying exposures.
- ▶ The outcome of a model with an offset is now  $\frac{y}{t}$ .

---

<sup>2</sup>The same logic would apply if we measured quantities over varying spatial units, e.g. counting people in blocks versus census tracts.



# Offsets

## Explanation

- ▶ The mean of a Poisson process,  $\lambda$  is implicitly  $\lambda = \frac{\mu}{\tau}$ , the expected number of events,  $\mu$ , over the duration  $\tau$ .
- ▶ Assume a Poisson process where  $\lambda_i$  is the expected number of events for the  $i^{th}$  observation. We can write the link function as

$$y = \text{Poisson}(\lambda)$$

$$\log(\lambda) = \log\left(\frac{\mu}{\tau}\right) = \beta_0 + \beta_1 x$$

- ▶ This can be re-written as

$$= \log(\mu) - \log(\tau) = \beta_0 + \beta_1 x$$

# Offsets

## Explanation

- ▶ We can think of  $\tau$  as the number of **exposures** for each observation. Thus, we can write out a new model for  $\mu$ :

$$y \sim \text{Poisson}(\mu)$$

$$\log(\mu) = \log(\tau) + \beta_0 + \beta_1 x$$

# Offsets

## Simulated example

```
N <- 1000
tweets <- sample(c(1:100), N, replace = TRUE)
ideology <- rbinom(N,1,0.4)
likes <- c()
for (i in 1:N) {
  y <- sum(rpois(tweets[i], exp(1 + 1*ideology + rnorm(1))))
  likes[i] <- y
}
sims <- as_tibble(cbind(tweets, likes, ideology))
```

# Offsets

## Simulated example

```
head(sims)
```

```
## # A tibble: 6 x 3
##   tweets likes ideology
##   <int> <int>   <int>
## 1     15   231       0
## 2     15    26       0
## 3      9    30       1
## 4     48   109       1
## 5      6    46       0
## 6     47    93       1
```

# Offsets

## Specification and interpretation

- ▶ The model is specified by adding the logarithm of exposures (e.g.  $\log(\tau)$ ) as an **offset** using the `offset` function.
  - ▶ The coefficient for the logarithm of exposures is fixed to  $\beta_{\text{offset}} = 1$ .
- ▶ The model is now interpreted as predicting a *rate* rather than a count.
- ▶ We could also directly include the logarithm of exposures as a predictor and let the model determine the coefficient.

# Offsets

## Simulated example

```
m1 <- glm(likes ~ 1 + ideology,  
          data = sims, family = poisson(link = "log"))  
m2 <- glm(likes ~ ideology + log(tweets),  
          data = sims, family = poisson(link = "log"))  
m3 <- glm(likes ~ ideology + offset(log(tweets)),  
          data = sims, family = poisson(link = "log"))
```

# Offsets

	Poisson	Poisson (Log exposure)	Poisson (Offset)
(Intercept)	5.888*** (0.002)	2.073*** (0.014)	1.989*** (0.002)
ideology	-0.150*** (0.004)	-0.175*** (0.004)	-0.176*** (0.004)
log(tweets)		0.980*** (0.003)	
Num.Obs.	1000	1000	1000
Log.Lik.	-231359.705	-165445.532	-165464.738
F	1738.806	45275.265	2388.381

+  $p < 0.1$ , \*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

# Offsets

	Poisson	Poisson (Log exposure)	Poisson (Offset)
(Intercept)	360.718*** (0.002)	7.947*** (0.014)	7.309*** (0.002)
ideology	0.861*** (0.004)	0.839*** (0.004)	0.839*** (0.004)
log(tweets)		2.663*** (0.003)	
Num.Obs.	1000	1000	1000
Log.Lik.	-231359.705	-165445.532	-165464.738
F	1738.806	45275.265	2388.381

Exponentiated coefficients.

+  $p < 0.1$ , \*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$



# Offsets

## Example: Predicting retweet rates

- ▶ Three models of yearly retweets
  - ▶ No offset
  - ▶  $\text{Log}(\text{tweets})$  included as predictor
  - ▶  $\text{Log}(\text{tweets})$  included as offset

# Offsets

## Example: Predicting retweet rates

```
nb.rt <- glm.nb(retweet_total ~  
                populist + left_right + seats_per + country,  
                data = data)  
nb.rt.e <- glm.nb(retweet_total ~ log(tweet_count) +  
                  populist + left_right + seats_per + country,  
                  data = data)  
nb.rt.o <- glm.nb(retweet_total ~ offset(log(tweet_count)) +  
                  populist + left_right + seats_per + country,  
                  data = data)
```

# Offsets

	NB	NB (Log exposure)	NB (Offset)
(Intercept)	9.795*** (0.599)	0.336 (0.561)	2.570*** (0.456)
populist	0.375 (0.255)	0.068 (0.189)	0.163 (0.194)
left_right	-0.047 (0.041)	0.000 (0.030)	-0.017 (0.031)
seats_per	0.039*** (0.008)	0.020*** (0.006)	0.026*** (0.006)
log(tweet_count)		1.324*** (0.053)	
Num.Obs.	255	255	255
Log.Lik.	-2858.070	-2752.504	-2762.354
F	25.315	56.445	13.570

+  $p < 0.1$ , \*  $p < 0.05$ , \*\*  $p < 0.01$ , \*\*\*  $p < 0.001$

# Offsets

## Using offsets

- ▶ Include an offset if there are differences in measurement intervals across observations.
- ▶ Offsets allow models to be interpreted as rates rather than counts.
- ▶ The logarithm of exposures can also be directly modeled, but interpretation is less intuitive.

# Zero-inflated models

## Intuition

- ▶ Some count outcomes have high rates of zeros. What if the outcomes with a value of zero are generated by a different kind of process?
- ▶ **Zero-inflated models** allow us to separately model the process determining whether counts are non-zero and the expected count for each observations.

# Zero-inflated models

## Specification

- ▶ The zero-inflated Poisson model consists of a mixture of two linear models, a logistic regression predicting the probability of a zero and a Poisson model predicting the count outcome.

$$y_i = ZIPoisson(p, \lambda)$$

$$\text{logit}(p) = \beta_{0p} + \beta_{1p}x$$

$$\log(\lambda) = \beta_{0\lambda} + \beta_{1\lambda}x$$

- ▶ Each model has its own parameters. These can be specified to model each process.

# Zero-inflated models

## Example: Books borrowed from the library

```
N <- 100
prob_lib <- 0.6
lib <- rbinom(N, 1, prob_lib)
sum(lib)/N
## [1] 0.47
```

# Zero-inflated models

## Example: Books borrowed from the library

```
x <- rnorm(N)

books <- c()
for (i in 1:N) {
  if (lib[i] == 1) {
    b <- rpois(1, lambda = exp(1 + 0.3*x[i] + rnorm(1)))
    books[i] <- b
  }
  else {books[i] <- 0}
}
mean(books)

## [1] 2.37

max(books)

## [1] 62
```



# Zero-inflated models

## Two kinds of zeros

```
sum(books == 0)
```

```
## [1] 65
```


```
sum(books == 0 & lib == 1)
```

```
## [1] 12
```

```
sum(books == 0 & lib == 0)
```

```
## [1] 53
```

# Zero-inflated models

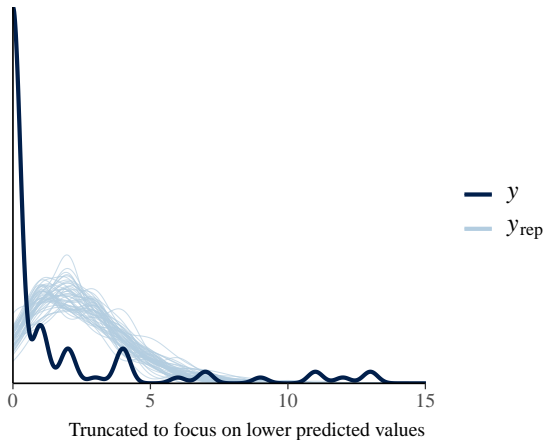


# Zero-inflated models

## Estimating a Poisson model


```
book.data <- as.data.frame(cbind(books, x))  
pois.m <- stan_glm(books ~ x, data = book.data, family = poisson(),  
  seed = 08901, chains = 1, refresh = 0)
```

# Zero-inflated models



# Zero-inflated models

Predicted number of zeros



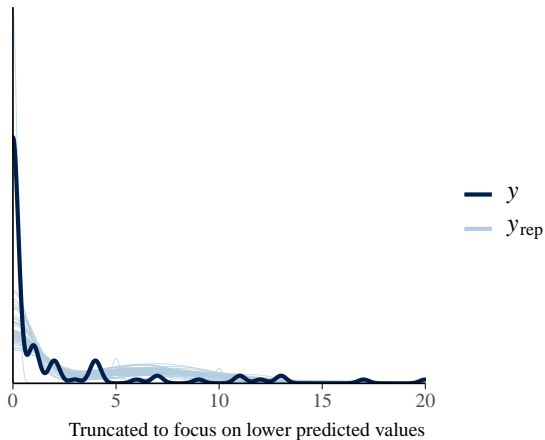
# Zero-inflated models

## Estimating a zero-inflated Poisson model

We must use the brms library to implement Bayesian zero-inflated Poisson regression.


```
library(brms)
zip <- brm(books ~ x,
          data = book.data,
          family = zero_inflated_poisson(link = "log",
                                         link_zi = "logit"),
          seed = 08901, refresh = 0, chains = 1)
```

# Zero-inflated models



# Zero-inflated models

Predicted number of zeros





# Zero-inflated models

## Comparing standard and zero-inflated models

##	elpd_diff	se_diff
## zip	0.0	0.0
## pois.m	-175.8	75.2

# Summary

- ▶ Standard linear models are generally unsuitable for count data
- ▶ Poisson regression can be used for most count outcomes
- ▶ Overdispersion occurs when variation higher than expected under Poisson model
  - ▶ Negative binomial regression includes a scale parameter
- ▶ Zero-inflated models are used to decompose processes generating zeros and counts

# Next week

- ▶ Categorical outcomes
  - ▶ Multinomial and ordered logistic regression